

(Downloaded from <https://cs.stanford.edu/~knuth/programs.html> and typeset on September 17, 2017)

1. Intro. This program was written (somewhat hastily) in order to experiment with an algorithm that generates all spanning trees of a given graph, changing only one edge at a time. Most of the basic ideas are adapted from Malcolm Smith’s M.S. thesis, “Generating spanning trees” (University of Victoria, 1997), which also contains more complex variations that guarantee better asymptotic performance. I intend to experiment with those additional bells and whistles later.

The first command line argument is the name of a file that specifies an undirected graph in Stanford GraphBase SAVE.GRAPH format; the graph may have repeated edges, but it must not contain loops. Additional command line arguments are ignored except that they cause more verbose output. The least verbose output contains only overall statistics about the total number of spanning trees found and the total number of mems used.

```
#define verbose (argc > 2)
#define extraverbose (argc > 3)
#define o mems++
#define oo mems += 2
#define ooo mems += 3
#define oooo mems += 4
#define ooooo mems += 5
#include "gb_graph.h"
#include "gb_save.h"
  ⟨Preprocessor definitions⟩
double mems; /* memory references made */
double count; /* trees found */
  ⟨Subroutines 5⟩
main(int argc, char *argv[])
{
  ⟨Local variables 3⟩;
  ⟨Input the graph 2⟩;
  ⟨Initialize the algorithm 16⟩;
  ⟨Generate all spanning trees 7⟩;
  printf("Altogether %.15g spanning trees, using %.15g mems.\n", count, mems);
  exit(0);
}
```

```
2.  ⟨Input the graph 2⟩ ≡
if (argc < 2) {
  fprintf(stderr, "Usage: %s foo.gb [[gory] details]\n", argv[0]);
  exit(1);
}
g = restore_graph(argv[1]);
if (!g) {
  fprintf(stderr, "Sorry, can't create the graph from file %s! (error code %d)\n", argv[1],
    panic_code);
  exit(-1);
}
n = g-n;
  ⟨Check the graph for validity and prepare it for action 4⟩;
```

This code is used in section 1.

3. $\langle \text{Local variables 3} \rangle \equiv$
register Graph **g*; /* the graph we're dealing with */
register int *n*; /* the number of vertices */
register int *k*; /* current integer of interest */
register Vertex **u, *v, *w*; /* current vertices of interest */
register Arc **e, *ee, *f, *ff*; /* current edges of interest */

See also section 8.

This code is used in section 1.

4. Graph preparation. While we’re checking to see that the graph meets certain minimal standards, we might as well also compute the degree of each vertex, since our algorithm will be using that information. We also ensure that the SGB “edge trick” works on our computer.

In this program we deviate from normal conventions of the Stanford GraphBase by using a *doubly* linked list of arcs from each vertex v . Namely, v -arcs points to a header node h , and the arcs from v are h -next, h -next-next, etc., until returning to h again. All arc nodes e in this list have e -next-prev = e -prev-next = e . The header node is distinguished by the property h -tip = Λ .

The “length” of each edge is changed to an identifying number for use in printouts.

```
#define deg u.I /* utility field u of each vertex holds its current degree */
#define prev a.A /* utility field a of each arc holds its backpointer */
#define mate(e) (edge_trick & (siz_t)(e) ? (e) - 1 : (e) + 1)

< Check the graph for validity and prepare it for action 4 > ≡
if (verbose) printf("Graph%s has the following edges:\n", g-id);
for (v = g-vertices, k = 0; v < g-vertices + n; v++) {
    f = gb_virgin_arc();
    f-next = v-arcs; /* the new header node */
    for (v-deg = 0, e = v-arcs, v-arcs = f; e; v-deg++, f = e, e = e-next) {
        e-prev = f;
        u = e-tip;
        if (u ≡ v) {
            fprintf(stderr, "Oops, there's a loop from %s to itself!\n", v-name);
            exit(-3);
        }
        if (mate(e)-tip ≠ v) {
            fprintf(stderr, "Oops: There's an arc from %s to %s,\n", u-name, v-name);
            fprintf(stderr, "but the edge_trick doesn't find the opposite arc!\n");
            exit(-4);
        }
        if (u > v) {
            e-len = mate(e)-len = ++k;
            if (verbose) printf("%d: %s--%s\n", k, v-name, u-name);
        }
    }
    v-arcs-prev = f, f-next = v-arcs; /* complete the double linking */
    if (v-deg ≡ 0) {
        fprintf(stderr, "Graph%s has an isolated vertex!\n", g-id, v-name);
        exit(-5);
    }
}
```

This code is used in section 2.

5. Here’s something I might like to use when debugging.

```
< Subroutines 5 > ≡
void print_arcs(Vertex *v)
{
    register Arc *a;
    printf("Arcs leading from %s:\n", v-name);
    for (a = v-arcs-next; a-tip; a = a-next) printf("%d(to %s)\n", a-len, a-tip-name);
}
```

See also section 17.

This code is used in section 1.

6. The method. Let G be a graph with $n > 1$ vertices. The basic idea of Smith's algorithm is to generate all spanning trees of G in such a way that the first one includes a given *near-tree*, namely a set of $n - 2$ edges that don't contain a cycle. This task is easy if $n = 2$: We simply list all the edges.

If $n > 2$ and if the near-tree is $\{e_1, \dots, e_{n-2}\}$, we proceed as follows: First form the graph $G \cdot e_1$ by shrinking edge e_1 (making its endpoints identical). All spanning trees of G that include e_1 are obtained by appending e_1 to a spanning tree of $G \cdot e_1$; so we proceed recursively to generate all spanning trees of $G \cdot e_1$, beginning with the near-tree $\{e_2, \dots, e_{n-2}\}$. If no such trees exist, we stop; in this case $G \cdot e_1$ is not connected, so G is not connected and it has no spanning trees. Otherwise, suppose the last spanning tree found for $G \cdot e_1$ is $f_1 \dots f_{n-2}$. Then we complete our task by deleting edge e_1 and generating all spanning trees in the resulting graph $G \setminus e_1$, starting with the near-tree $\{f_1, \dots, f_{n-2}\}$.

7. This program implements the recursion directly by maintaining an array of edges $a_1 \dots a_n$. When we enter level l , positions $a_1 \dots a_{l-1}$ contain edges to include in the tree, and those edges have been shrunk in the current graph. Position a_l is effectively blank, and the remaining positions $a_{l+1} \dots a_n$ contain the edges of a near-tree that should be part of the next spanning tree generated.

We don't delete an edge that is a "bridge," whose removal would disconnect the current graph. When a non-bridge edge e is deleted at level l , we set $change_e = e$. If the previously found spanning tree was $a_1 \dots a_{n-1}$, the next tree found will be $a_1 \dots a_{l-1} a_{l+1} \dots a_{n-1} e'$ for some new edge e' ; thus it will differ from its predecessor by removing edge $change_e$ and replacing it with e' .

It's convenient to keep array element a_l in a utility field within the vertex array, represented by $aa(l)$. Another such utility field, $del(l)$, points to a stack of the edges deleted before coming to a bridge; edges on this list are linked together via a *link* field.

Experienced readers will not be shocked by the fact that this part of the program has a **goto** leading from one loop into another.

```
#define aa(l) (g-vertices + l)-z.A    /* the edge a_l */
#define del(l) (g-vertices + l)-y.A    /* the most recent edge deleted on level l */
#define link b.A    /* points from one edge to another */
⟨Generate all spanning trees 7⟩ ≡
    change_e = Λ;
    v = g-vertices;    /* this instruction needed only if n = 2 */
    for (l = 1; l < n - 1; l++) {
        o, del(l) = Λ;
        enter: ooo, e = aa(l + 1), u = e-tip, v = mate(e)-tip;
        if (oo, u-deg > v-deg) v = u, e = mate(e), u = e-tip;
        ⟨Shrink e by changing u to v 10⟩;
        o, aa(l) = e;
    }
    for (o, e = v-arcs-next; o, e-tip; o, e = e-next) {
        o, aa(l) = e;
        ⟨Produce a new spanning tree by changing change_e to e 9⟩;
        change_e = e;
    }
    for (l--; l; l--) {
        e = aa(l), u = e-tip, v = mate(e)-tip;
        ⟨Unshrink e by restoring u 11⟩;
        ⟨If e is not a bridge, delete it, set change_e = e, and goto enter 12⟩;
        ⟨Undelete all edges deleted since entering level l 15⟩;
    }
```

This code is used in section 1.

8. $\langle \text{Local variables } 3 \rangle + \equiv$
register **int** *l*; /* the current level */
Arc **change_e*; /* edge that may change next */

9. $\langle \text{Produce a new spanning tree by changing } change_e \text{ to } e \ 9 \rangle \equiv$
count++;
if (*verbose*) {
 if ($\neg change_e \vee extravertbose$) {
 printf("%15g:", *count*);
 for (*k* = 1; *k* < *n*; *k*++) *printf*("_%d", *aa*(*k*)-*len*);
 if (*extravertbose* \wedge *change_e*) *printf*("_(-%d+%d)\n", *change_e*-*len*, *e*-*len*);
 else *printf*("\n");
 } **else** *printf*("%15g:_-%d+%d\n", *count*, *change_e*-*len*, *e*-*len*);
} }

This code is used in section 7.

10. To shrink an edge between *u* and *v*, we insert *u*'s adjacency list into *v*'s, changing all references to *u* into references to *v*; those references occur in *tip* fields of mates of the arcs in *u*'s list.

We also delete all former edges between *u* and *v*, since those would otherwise become loops. Those former edges are linked together via their *link* fields, so that we can restore them later.

Note that *e*-*tip* = *u*, so *e* appears in the *v* list while *mate*(*e*) appears in the *u* list.

#define *delete*(*e*) *ee* = *e*, *oooo*, *ee*-*prev*-*next* = *ee*-*next*, *ee*-*next*-*prev* = *ee*-*prev*

$\langle \text{Shrink } e \text{ by changing } u \text{ to } v \ 10 \rangle \equiv$
oo, *k* = *u*-*deg* + *v*-*deg*;
for (*o*, *f* = *u*-*arcs*-*next*, *ff* = Λ ; *o*, *f*-*tip*; *o*, *f* = *f*-*next*)
 if (*f*-*tip* \equiv *v*) *delete*(*f*), *delete*(*mate*(*f*)), *k* -= 2, *o*, *f*-*link* = *ff*, *ff* = *f*;
 else *o*, *mate*(*f*)-*tip* = *v*;
oo, *e*-*link* = *ff*, *v*-*deg* = *k*;
if (*extravertbose*)
 printf("level_%d:_Shrinking_%d;_now_%s_has_degree_%d\n", *l*, *e*-*len*, *v*-*name*, *v*-*deg*);
o, *ff* = *v*-*arcs*; /* now *f* = *u*-*arcs*; we will merge the two lists */
oooo, *f*-*prev*-*next* = *ff*-*next*, *ff*-*next*-*prev* = *f*-*prev*;
ooo, *f*-*next*-*prev* = *ff*, *ff*-*next* = *f*-*next*;

This code is used in section 7.

11. Unshrinking uses the principle of “dancing links,” whereby we exploit the fact that previously deleted nodes still have good information in their *prev* and *next* fields, provided that we undelete in reverse order.

#define *undelete*(*e*) *ee* = *e*, *oooo*, *ee*-*next*-*prev* = *ee*, *ee*-*prev*-*next* = *ee*

$\langle \text{Unshrink } e \text{ by restoring } u \ 11 \rangle \equiv$
oo, *f* = *u*-*arcs*, *ff* = *v*-*arcs*;
ooo, *ff*-*next* = *f*-*prev*-*next*;
o, *ff*-*next*-*prev* = *ff*;
ooo, *f*-*prev*-*next* = *f*, *f*-*next*-*prev* = *f*;
for (*f* = *f*-*prev*; *o*, *f*-*tip*; *o*, *f* = *f*-*prev*) *o*, *mate*(*f*)-*tip* = *u*;
for (*oo*, *f* = *e*-*link*, *k* = *v*-*deg*; *f*; *o*, *f* = *f*-*link*) *k* += 2, *undelete*(*mate*(*f*)), *undelete*(*f*);
oo, *v*-*deg* = *k* - *u*-*deg*;
if (*extravertbose*)
 printf("level_%d:_Unshrinking_%d;_now_%s_has_degree_%d\n", *l*, *e*-*len*, *v*-*name*, *v*-*deg*);

This code is used in section 7.

12. For bridge detection, we try a heuristic that often gives a quick answer when the graph is sparse (namely, to test if u has degree 1). Or if $e\text{-link-link} \neq \Lambda$, there was another edge between u and v . Otherwise we resort to brute-force breadth-first, testing whether one can get from u to v without e .

When I put this algorithm in a book, I'll probably leave out the two quick-try heuristics, in order to keep the algorithm shorter; breadth-first search will resolve both cases without too much additional calculation. But for now I'm trying to see how useful they are.

```
#define bfs v.V /* link for the breadth-first search: nonnull if vertex seen */
⟨If  $e$  is not a bridge, delete it, set  $change\_e = e$ , and goto enter 12⟩ ≡
  if ( $o, u\text{-deg} \equiv 1$ ) {
    if (extraverbose) printf("level_%d:_%d_is_a_bridge_with_endpoint_%s\n", l, e-len, u-name);
    goto bridge;
  }
  if ( $o, e\text{-link-link}$ ) {
    if (extraverbose) printf("level_%d:_%d_is_parallel_to_%d\n", l, e-len,
      e-link-len ≠ e-len ? e-link-len : e-link-link-len);
    goto nonbridge;
  }
  for ( $o, u\text{-bfs} = v, w = u; u \neq v; o, u = u\text{-bfs}$ ) {
    for ( $oo, f = u\text{-arcs-next}; o, f\text{-tip}; o, f = f\text{-next}$ )
      if ( $o, f\text{-tip-bfs} \equiv \Lambda$ ) {
        if ( $f\text{-tip} \equiv v$ ) {
          if ( $f \neq mate(e)$ ) ⟨Nullify all bfs links and goto nonbridge 13⟩;
        } else  $oo, f\text{-tip-bfs} = v, w\text{-bfs} = f\text{-tip}, w = f\text{-tip};$ 
      }
  }
  if (extraverbose) printf("level_%d:_%d_is_a_bridge\n", l, e-len);
  for ( $o, u = e\text{-tip}; u \neq v; o, u\text{-bfs} = \Lambda, u = w$ )  $o, w = u\text{-bfs};$ 
  goto bridge;
nonbridge:  $change\_e = e;$ 
  ⟨Delete  $e$  and goto enter 14⟩;
bridge:
```

This code is used in section 7.

```
13. ⟨Nullify all bfs links and goto nonbridge 13⟩ ≡
{
  for ( $o, u = e\text{-tip}; u \neq v; o, u\text{-bfs} = \Lambda, u = w$ )  $o, w = u\text{-bfs};$ 
  goto nonbridge;
}
```

This code is used in section 12.

```
14. ⟨Delete  $e$  and goto enter 14⟩ ≡
  if (extraverbose) printf("level_%d:_%d_deleting_%d\n", l, e-len);
   $ooo, e\text{-link} = del(l), del(l) = e;$ 
  delete( $e$ ), delete( $mate(e)$ ),  $oo, e\text{-tip-deg} --, v\text{-deg} --;$ 
  goto enter;
```

This code is used in section 12.

15. \langle Undelete all edges deleted since entering level l 15 $\rangle \equiv$
for ($o, e = del(l); e; o, e = e\text{-}link$) {
 $oooo, mate(e)\text{-}tip\text{-}deg ++, e\text{-}tip\text{-}deg ++, undelete(mate(e)), undelete(e);$
 if ($extraverbose$) $printf("undeleting_d\n", e\text{-}len);$
}

This code is used in section 7.

16. Getting started. We're done, except for one embarrassing detail: It is necessary to prime the pump by setting up the original near-tree $a_2 \dots a_{n-1}$. For this purpose I'll use depth-first search, since it seems a bit faster than the alternatives. And I might as well check that the graph is connected.

```
#define sentinel (g-vertices)
⟨ Initialize the algorithm 16 ⟩ ≡
    for (v = g-vertices + 1; v < g-vertices + n; v++) v-bfs = Λ;
    for (k = n - 1, o, w = v = g-vertices, w-bfs = sentinel; ; o, v = w, w = w-bfs) {
        for (oo, e = v-arcs-next; o, u = e-tip; o, e = e-next)
            if (o, u-bfs ≡ Λ) {
                o, aa(k) = e, k--;
                if (k ≡ 0) goto connected;
                o, u-bfs = w, w = u;
            }
        if (w ≡ sentinel) break;
    }
    printf("Oops, the graph isn't connected!\n"); exit(0);
connected:
    for (u = g-vertices; u < g-vertices + n; u++) o, u-bfs = Λ;
    if (extraverbose) {
        printf("Depth-first search yields the following spanning tree:\n");
        print_a(g);
    }
    if (verbose) printf("(%.15g mems for initialization)\n", mems);
```

This code is used in section 1.

17. One final debugging aid.

```
⟨ Subroutines 5 ⟩ +≡
    void print_a(register Graph *g)
    {
        register int k;
        for (k = 1; k < g-n; k++)
            printf("a%d=%d (%s--%s)\n", k, aa(k)-len, aa(k)-tip-name, mate(aa(k))-tip-name);
    }
```


18. Index.

a: [5](#).
aa: [7](#), [9](#), [16](#), [17](#).
Arc: [3](#), [5](#), [8](#).
arcs: [4](#), [5](#), [7](#), [10](#), [11](#), [12](#), [16](#).
argc: [1](#), [2](#).
argv: [1](#), [2](#).
bfs: [12](#), [13](#), [16](#).
bridge: [12](#).
change_e: [7](#), [8](#), [9](#), [12](#).
connected: [16](#).
count: [1](#), [9](#).
deg: [4](#), [7](#), [10](#), [11](#), [12](#), [14](#), [15](#).
del: [7](#), [14](#), [15](#).
delete: [10](#), [14](#).
e: [3](#).
edge_trick: [4](#).
ee: [3](#), [10](#), [11](#).
enter: [7](#), [14](#).
exit: [1](#), [2](#), [4](#), [16](#).
extraverbose: [1](#), [9](#), [10](#), [11](#), [12](#), [14](#), [15](#), [16](#).
f: [3](#).
ff: [3](#), [10](#), [11](#).
fprintf: [2](#), [4](#).
g: [3](#), [17](#).
gb_virgin_arc: [4](#).
Graph: [3](#), [17](#).
id: [4](#).
k: [3](#), [17](#).
l: [8](#).
len: [4](#), [5](#), [9](#), [10](#), [11](#), [12](#), [14](#), [15](#), [17](#).
link: [7](#), [10](#), [11](#), [12](#), [14](#), [15](#).
main: [1](#).
mate: [4](#), [7](#), [10](#), [11](#), [12](#), [14](#), [15](#), [17](#).
mems: [1](#), [16](#).
n: [3](#).
name: [4](#), [5](#), [10](#), [11](#), [12](#), [17](#).
next: [4](#), [5](#), [7](#), [10](#), [11](#), [12](#), [16](#).
nonbridge: [12](#), [13](#).
o: [1](#).
oo: [1](#), [7](#), [10](#), [11](#), [12](#), [14](#), [16](#).
ooo: [1](#), [7](#), [10](#), [11](#), [14](#).
oooo: [1](#), [10](#), [11](#), [15](#).
ooooo: [1](#).
panic_code: [2](#).
prev: [4](#), [10](#), [11](#).
print_a: [16](#), [17](#).
print_arcs: [5](#).
printf: [1](#), [4](#), [5](#), [9](#), [10](#), [11](#), [12](#), [14](#), [15](#), [16](#), [17](#).
restore_graph: [2](#).
sentinel: [16](#).
siz_t: [4](#).
stderr: [2](#), [4](#).
tip: [4](#), [5](#), [7](#), [10](#), [11](#), [12](#), [13](#), [14](#), [15](#), [16](#), [17](#).
u: [3](#).
undelele: [11](#), [15](#).
v: [3](#), [5](#).
verbose: [1](#), [4](#), [9](#), [16](#).
Vertex: [3](#), [5](#).
vertices: [4](#), [7](#), [16](#).
w: [3](#).

- ⟨ Check the graph for validity and prepare it for action 4 ⟩ Used in section 2.
- ⟨ Delete e and **goto** *enter* 14 ⟩ Used in section 12.
- ⟨ Generate all spanning trees 7 ⟩ Used in section 1.
- ⟨ If e is not a bridge, delete it, set $change_e = e$, and **goto** *enter* 12 ⟩ Used in section 7.
- ⟨ Initialize the algorithm 16 ⟩ Used in section 1.
- ⟨ Input the graph 2 ⟩ Used in section 1.
- ⟨ Local variables 3, 8 ⟩ Used in section 1.
- ⟨ Nullify all *bfs* links and **goto** *nonbridge* 13 ⟩ Used in section 12.
- ⟨ Produce a new spanning tree by changing $change_e$ to e 9 ⟩ Used in section 7.
- ⟨ Shrink e by changing u to v 10 ⟩ Used in section 7.
- ⟨ Subroutines 5, 17 ⟩ Used in section 1.
- ⟨ Undelete all edges deleted since entering level l 15 ⟩ Used in section 7.
- ⟨ Unshrink e by restoring u 11 ⟩ Used in section 7.

GRAYSPAN

	Section	Page
Intro	1	1
Graph preparation	4	3
The method	6	4
Getting started	16	8
Index	18	9